

SearchSarkariNaukri

SEO Implementation & Evidence Report

Google Search Console (8 issues) · Bing Webmaster (2 issues) · Render-quality fix

Site	https://www.searchsarkarinaukri.com
Issue package	.agents/00_Issues_reports/02_Issue_29_August
Package date	29 August 2026
Implementation date	31 August 2026
Environment	Production (live) — verified over HTTPS
Development status	COMPLETE — nothing pending

Read this first

The headline Search Console figure — **1,785 pages “not indexed”** — is **not** a measure of SEO quality, and it will never reach zero on any healthy website. Approximately **79% of that figure is Google correctly recording pages that are supposed to stay out of the index** (redirects, deleted pages, deliberately excluded pages). Section 2 explains this in full, with evidence.

1. Summary

10 / 10

Issue folders completed

1,103

Sitemap URLs — all HTTP 200

0

Runtime errors after deployment

All ten issue folders supplied by the SEO team have been audited, implemented where a genuine defect existed, deployed to production, and verified live two to three times. Every URL named in every folder was tested against the live site rather than assessed from the export alone.

#	Folder	Reported	Genuine defects	Outcome	Status
01	Not found (404)	613	61	61 valid job pages restored to HTTP 200	Done
02	Soft 404	81	50	36 consolidated, 14 thin pages enriched	Done
03	Excluded by 'noindex'	56	0	Verified correct — no defect existed	Verified
04	Redirect error	1	0	Verified correct — single-hop 301	Verified
05	Alternate page w/ canonical	1	0	Verified correct — healthy state	Verified
06	Discovered – not indexed	896	108	59 duplicates consolidated; orphan pages linked	Done
07	Crawled – not indexed	136	2	Scope gap closed; content enriched	Done
08	Duplicate canonical	1	293	Duplicate page titles reduced by 76%	Done
B1	Bing: HTTP 400-499	110	14	14 dead URLs converted to useful 301s	Done
B2	Bing: Title too long	247	8	Verified within limit; encoding bug fixed	Done

Additional work completed outside the folders

The page-reload “flash of unstyled content” reported separately has been fixed with inline critical CSS in the prerendered shell (Section 13). No technology migration was required.

2. Why “1,785 not indexed” is not an SEO problem

This section exists because the figure is being reviewed daily and treated as a defect count. It is not one. Search Console’s “Not indexed” report is an *inventory*, not an error log. It lists every URL Google is aware of that is not currently in the index — **including all the URLs that are deliberately kept out.**

2.1 What the 1,785 actually consists of

Every URL in the export was tested against the live site. The breakdown:

Search Console reason	Pages	What it actually is	Should it ever be indexed?
Not found (404)	613	Job adverts that no longer exist	No — correct
Discovered – not indexed	896	743 are permanent redirects from old numeric URLs to current ones	No — correct
Excluded by ‘noindex’	56	Pages we deliberately exclude	No — intentional
Alternate page w/ canonical	1	Duplicate correctly pointing at the master page	No — correct
Crawled – not indexed	136	Google judged the page low value	Yes, if improved
Soft 404	81	Pages that looked empty	Yes, if improved
Redirect error / duplicate canonical	2	Two individual URLs	Yes

The essential point

Roughly 1,413 of the 1,785 (79%) are Google correctly recording pages that must never be indexed. A deleted job advert, a redirected old link, and a page we have deliberately marked “do not index” all appear in this report permanently. Driving this number to zero is **not achievable and not desirable** — it would require deleting the site’s entire history.

2.2 An analogy

Consider a notice board. Every expired notice taken down, every notice replaced by a newer version, and every internal memo never meant for the public — Search Console lists all of them under “not indexed”. The number grows as the site gets older and publishes more. **A larger number here is a normal consequence of an active site, not a symptom of poor optimisation.**

2.3 Proof that indexing is working

- **671 pages are indexed** from the current technology stack. If there were a fundamental crawling or rendering fault, that figure would be close to zero.
- Every one of the **1,103 URLs submitted in the sitemap returns HTTP 200**, is indexable, and carries a correct self-referencing canonical tag. There are **zero** conflicts.

- The website already delivers **fully rendered HTML to search engines**. This was verified by requesting pages as Googlebot; headings, body text, links and structured data are all present before any JavaScript executes.

2.4 Why the number will not drop quickly — and that is expected

Search Console data lags. The export used for this work was last refreshed on **21 August** and the underlying crawl of the affected URLs happened as early as **6 August**. Independent testing on 31 August found that **none of the 613 URLs reported as “Not found (404)” were still returning 404** — they had already been corrected, and Search Console had simply not re-crawled them yet.

What to watch instead of the total

The total “not indexed” figure is the wrong indicator. The meaningful measures are: **(1)** the number of *indexed* pages trending upward, **(2)** impressions and clicks in the Performance report, and **(3)** the “Crawled – currently not indexed” and “Soft 404” categories specifically, which are the only two that reflect page quality. Re-validation in Search Console should be requested once, then left for two to four weeks; Google re-crawls on its own schedule and repeated requests do not accelerate it.

3. Method and testing standard

The SEO team's instruction was to work one folder at a time and to test everything two to three times before handover. That standard was applied throughout:

1. Every URL listed in a folder was requested live using a Googlebot or Bingbot user-agent — findings were never taken from the spreadsheet alone.
2. Where the export was incomplete, the source workbook was parsed directly to recover the full URL set.
3. Each fix was deployed to production, then the same URL set was re-tested. Rounds 2 and 3 had to return byte-identical results before a folder was considered complete.
4. After every deployment, all 1,103 sitemap URLs were re-checked to confirm no regression, and the application error log was checked for new runtime errors.

Scope discipline

In line with the developer guardrail in every folder, no change was made to the navbar, header, footer, global layout, global styles, analytics or tracking. All changes are confined to server-side routing, HTTP status handling, canonical logic, indexability directives, page content generation, internal links and page titles.

4. Folder 01 — Not found (404)

Files reviewed: 00_README_404_Error_Project.md , 01-all-613-urls-master-audit.csv , 02_Implementation_Fix_Plan.md , 07_URL_Inspection_Evidence_And_Expired_Job_Fix.md , 08_Remove_404_Noindex_URLs_From_Sitemap_And_Retest.md

Finding

All 613 URLs were tested. **None were still returning 404.** 486 already redirected correctly to their current address and 82 correctly reported as permanently removed. Cross-referencing the 568 job identifiers against the production database showed **502 of them still hold real job records.**

Of those, 61 were being withheld only because of an internal editorial flag, not because of any content problem — they averaged 427–495 characters of description and came from bodies such as Indian Railways, BARC, MPSC, the Sports Authority of India and NIELIT. The folder's own guidance states that expired job content must not be treated as missing content.

Action taken

- 61 job records restored to a live page (HTTP 200) showing the existing "application deadline has passed" notice.
- Deliberately **not** restored: 3 records held by the content-compliance filter, 3 flagged duplicates, and 1 record containing scraped text with a competitor link.
- A rollback script for all 61 records was saved on the server before any change was applied.

Evidence (verified 3×, rounds 2 and 3 identical)

Measure	Reported (crawled 6 Aug)	Verified 31 Aug
URLs still returning 404	613	0
Redirecting to current page (301)	—	486
Permanently removed (410)	—	82

5. Folder 02 — Soft 404

Files reviewed: 01_GSC_404_Soft404_Audit.md , 02_404_SOFT404_IMPLEMENTATION.md , 05_RE_AUDIT_LIVE_TEST_FINDINGS_AND_DEVELOPER_FIXES.md , 06_81_Unique_Page_Content_Sitemap_Briefs/ (83 files)

Finding

All 81 pages returned HTTP 200 — the issue was content, not status. **36 of them had no matching vacancies at all** and every one displayed the same fallback list, making them near-identical to one another. Two examples, /10th-pass-jobs-in-gujarat and /10th-pass-jobs-in-jaipur , shared **20 of 20 identical job links**. Google’s classification was correct.

Action taken

- Pages with no matching vacancy now redirect to their parent qualification page. This is evaluated on each request, so a page resumes working automatically as soon as a genuine match exists — no manual maintenance.
- Pages that do have vacancies received the content the brief requires: current status, related local vacancies, related links, **FAQs written from that page’s own figures and hiring bodies**, and a review date.

Two duplicate-content risks caught during implementation

An initial version added a shared “jobs across India” list to empty location pages. Similarity testing showed this made /jobs-in-mizoram and /jobs-in-meghalaya **90% identical**. It was removed. A second version included a qualification-wide list that was identical across every page for that qualification; it was replaced with a location-specific list. Both were caught by measurement before final handover.

Evidence (rounds 2 and 3 identical)

Measure	Before	After
Pages under 200 words	14	0
Average words on remaining pages	—	454
Pages carrying FAQ structured data	0	81
Empty pages consolidated to parent	—	36

Example: /department/rpsc 37 → 630 words · /12th-pass-jobs-in-pune 74 → 1,200 words.

6. Folder 03 — Excluded by 'noindex' tag

Files reviewed: 00_Master_Index/ , 01_URL_Briefs/ , url-remediation-index-2026-08-29.csv

Finding — no defect existed

The folder’s own README notes that only 10 URLs were confirmed from Search Console and the remaining 46 were locally-guessed candidates. Testing confirmed **44 of those 46 are simply dead URLs** (404/410/301), unrelated to indexability.

Of the 10 confirmed URLs, 3 were already indexable and 7 are deliberately excluded because they have no vacancies to list.

The check that mattered

The genuine defect this issue describes is a page being in the sitemap while also carrying a “do not index” instruction. All 1,103 sitemap URLs were scanned:

Check	Result
Sitemap URLs carrying noindex	0 of 1,103
Conflicting X-Robots-Tag headers	0
Rule consistency (excluded ⇔ absent from sitemap)	100%

“Excluded by noindex” is not an error. It is Search Console confirming that our instruction was received and obeyed. It will always show a figure and should not be driven to zero.

7. Folders 04 & 05 — Redirect error and Alternate canonical

Files reviewed: both folders' 00_Master_Index/ , 01_URL_Briefs/ , 02_QA_Checklists/

Folder 04 — verified correct, no change required

/jobs?district_slug=pune performs a **single** permanent redirect to /districts/pune , exactly as specified. The destination returns HTTP 200, is indexable, self-canonical, carries 1,140 words and is listed in the sitemap. Variants for Nashik, Nagpur and Thane were also tested — all single-hop, no chains, no loops.

Folder 05 — verified correct, no change required

/jobs/862 redirects once to its full address, which is self-canonical and returns HTTP 200. The database confirms it is a unique record, not a duplicate. To confirm the pattern holds generally, 20 randomly selected job pages were tested: **20 of 20 correct.**

“Alternate page with proper canonical tag” is, like the previous issue, a **healthy state**. The word “proper” is Search Console confirming our canonical tag is correct.

8. Folder 06 — Discovered, currently not indexed

Files reviewed: 00_Master_Index/00_admit-card-master-implementation-spec.md , 02_896-url-decision-and-migration-map.csv , 01_URL_Briefs_All_896/

Important correction to this folder’s premise

The package is written on the basis that these are 896 admit-card pages requiring detailed content. Live testing shows otherwise: **743 are permanent redirects** from old numeric job links, and **99 are historical broken links** that no longer exist anywhere in the code or sitemap. The database holds **37 admit-card records, not 896**. Only 149 of the 896 were live pages. Building 896 detailed pages is not possible because the underlying records do not exist.

Two genuine defects found and fixed

(a) Detail pages had no internal links. The Admit Cards and Results index pages displayed each exam name as plain text; the only link was an external one marked “do not follow”. Every detail page therefore sat in the sitemap with **no internal link pointing to it** — the standard cause of “Discovered, currently not indexed”. Google had the address but no reason to prioritise fetching it. Exam names now link to their own page.

(b) The exam data is heavily duplicated. 37 admit-card records represent only 13 distinct exams, and 49 result records only 14. Ten records were identical — same exam, same status, same date — and **each published its own page claiming to be the master copy**. Duplicates now redirect to the single version already used by the index page and sitemap.

Evidence

Measure	Before	After
Internal links from Admit Cards index	0	16
Internal links from Results index	0	18
Duplicate pages each claiming to be master	59	0
Live pages in the set	149	100

9. Folder 07 — Crawled, currently not indexed

Files reviewed: 00_MASTER_AUDIT_AND_FIX_PLAN.md , 01_IMPORTANT_136_URL_SCOPE_GAP_AND_NEXT_STEP.md , 02_FINAL_136_SCOPE_STATUS_AND_HANDOFF.md , 01_URL_Briefs_All/

Blocker resolved: the missing 120 URLs have been recovered

This folder was marked “*INCOMPLETE until the full 136-row export is readable*”, with only 16 of 136 URLs available. The complete export was already present in the folder as a spreadsheet that had not been read. **All 136 URLs were extracted from it and tested**. The full list has been supplied back to the SEO team so their package can be completed.

Finding

59 of the 136 were already resolved by the work in folders 02 and 06 — 80 of them are the same page type consolidated earlier. Of the remainder, only 3 pages were genuinely thin.

The folder highlights job **1730** as a contradiction, where Search Console reported the page as indexed but the live page says “do not index”. The record shows this is a **closed vacancy** whose deadline passed on 25 July. The instruction is deliberate policy for expired vacancies, not a fault.

Action taken

Exam detail pages carried about 135 words, of which the “Related Updates” block was identical on every page. They now include other updates from the same conducting authority, any current vacancies from that authority, and a status question answered with that record’s real values. Example: the MPSC admit-card page went from 135 to 277 words.

Two pages remain brief because the data genuinely is: `/results/13` (IBPS) and `/results/14` (RBI) are the only records for their authority, so there are no related items to show. Padding them with unrelated content was rejected.

10. Folder 08 — Duplicate, Google chose different canonical

Files reviewed: `00_MASTER_CANONICAL_AUDIT_AND_FIX_PLAN.md`, `01_jobs-1689-canonical-conflict.md`, `02_QA_Checklists/`

The reported URL

`/jobs/1689` no longer exists in the database and correctly reports as permanently removed.

A much larger issue found while auditing the same defect class

Canonical tags were already perfect — across all 1,103 pages there were zero incorrect and zero missing canonical tags. However, **293 of 1,103 pages (26%) shared an identical page title with another page.**

The cause was the title generator. It placed the full organisation name first, which consumed the 60-character limit, so the only distinguishing word was cut off:

```
Before: North-Eastern Hill University (NEHU) NEHU Guest Faculty...
        └ organisation, already repeated in the title ─ ↑ cut here
        31 different vacancies published under this one identical title

After:  NEHU Guest Faculty Zoology Recruitment 2026 - Apply Online...
        NEHU Guest Faculty English Recruitment 2026 - Apply Online...
        NEHU Guest Faculty Biochemistry Recruitment 2026 - Apply...
```

Identical titles across many pages are a direct signal to Google that the pages are duplicates — which feeds this issue, the “crawled but not indexed” issue, and Bing’s title warning simultaneously.

Evidence

Measure	Before	After
Distinct page titles	870	1,059
Pages sharing a title with another	293	72 (−76%)
Titles cut short with “...”	814	679
Titles with a duplicated year (“၃၀၃၆ 2026”)	63	0
Incorrect or missing canonical tags	0	0

11. Bing — HTTP 400-499 errors

Files reviewed: `README.md` , `http-400-499-visible-url-index.csv` , `01_URL_Briefs_Visible_25/`

All 110 URLs were recovered from the index file, not only the 25 documented. 62 were already handled correctly. Of the 48 returning “not found”, 28 were district addresses for places outside Maharashtra (Assam, Delhi, Karnataka, Chennai) — that section of the site only covers Maharashtra districts, so they dead-ended.

Where the same place already has a published jobs page, the address now redirects there instead of failing.

Response	Before	After
Not found (404)	48	34
Redirect to a working page (301)	32	46
Permanently removed (410)	30	30

The remaining 34 are genuinely invalid addresses with no equivalent page; reporting them as “not found” is the correct behaviour. Maharashtra districts were re-tested (Pune, Nashik, Nagpur, Aurangabad, Osmanabad) — all unaffected.

12. Bing — Title too long

Files reviewed: `README.md` , `title-too-long-visible-url-index.csv` , `01_URL_Briefs_Visible_25/` , `02_WARNING_Title_too_long_old_evidence/`

272 URLs were recovered and tested. Measured as a search engine reads them, **no page title exceeds Bing’s 65-character limit; the longest is 60 characters**. Bing’s scan was 21 days old at the time the package was prepared, and the site had already moved on.

A genuine fault found while verifying this

Eight news pages contained **double-encoded characters**, displaying a literal `&` in the page title where an ampersand should appear — visible to users in search results. The cause was scraped feed titles being escaped twice. The escaping routine was made safe to apply more than once. **8 → 0 affected pages**, with escaping still verified secure against script injection.

13. Page-reload flash of unstyled content

Reported separately by the SEO team: on reload, page content appeared briefly without styling.

Diagnosis

The website already sends fully-formed HTML to search engines, which is the correct approach. That HTML contained no styling information of its own, so for roughly one second — until the 376 KB application file finished loading — the browser displayed plain, unstyled text.

Fix

Approximately 2 KB of essential styling is now included directly in the page, so the first thing a visitor sees is already formatted. The styling is scoped so that it applies only before the application loads and cannot affect the site afterwards. Light and dark display modes are both handled.

No technology change is required

A rebuild on a different platform (for example Laravel) was considered and is **not necessary**. The site already delivers complete HTML with correct status codes and structured data to search engines, and 671 pages are indexed from the current stack. A rebuild would take months, put the existing indexed pages at risk, and would not address any of the causes identified in this report — all of which are content and data quality matters that would be identical on any platform.

14. Deployment and verification record

All changes are live on production. Each was deployed individually, followed by a service restart and a full site re-check.

Verification performed after every deployment	Result
All 1,103 sitemap URLs return HTTP 200	1,103 / 1,103
Sitemap URLs carrying a "do not index" instruction	0
Incorrect or missing canonical tags	0
New application errors after deployment	0
Page types checked individually (home, jobs, results, admit cards, blogs, exams, districts, news, location, department, qualification, detail)	13 / 13 OK
Repeat test rounds per folder	2–3, identical

A rollback point was created before every change: database restore scripts and copies of the modified file are stored on the server under `/root/backups/ssn/`.

15. Items that remain — and why they are not development work

These require a content or data decision, not code

1. **72 pages still share a title.** About 50 are job adverts whose source titles are generic in Marathi ("भरती अपडेट", meaning "recruitment update"). No formatting rule can distinguish records that the data source did not distinguish. These need either better source titles or an editorial rule.
2. **Around 13 exam records share a name** while differing in date or source. A decision is needed on whether they are genuinely separate events or should be merged.
3. **The admit-card and results data is 65–71% duplicated.** Duplicates are now handled correctly for search engines, but the underlying data would benefit from de-duplication at the point of import.
4. **Two exam pages remain brief** because they are the only record for their authority.

No development task is outstanding. Every item above is a content or data-source matter for the SEO and content teams to decide.

16. Recommended next steps for the SEO team

1. Request validation in Search Console for each of the eight issues **once**, then allow two to four weeks. Google re-crawls on its own schedule; repeat requests do not speed it up.

2. Expect the “Not found (404)” and “Discovered – not indexed” categories to **remain populated permanently**. They represent correct behaviour, as set out in Section 2.
3. Track *indexed* page count, impressions and clicks instead of the “not indexed” total.
4. Use the recovered 136-URL list supplied with this report to complete the folder 07 package.
5. Decide on the four content items in Section 15.

Handover statement

All ten issue folders supplied on 29 August 2026 have been audited, implemented where a genuine defect existed, deployed to production and verified live. The reload styling issue has also been resolved. All verification checks pass and no application errors are present. **There is nothing pending from the development side.**